Phase One Brief — Executive Summary

This summary is written for listening. Sentences are short. Acronyms are spelled out the first time they appear. There are no tables or code.

What Stewardly is

Stewardly is infrastructure. It supports all artificial-intelligence tools, all artificial-intelligence agents, and all related tooling that users want to bring with them. That is the foundational frame.

The platform also ships with a great out-of-box experience. A user who arrives without their own artificial-intelligence provider gets immediate value from the bundled defaults. A user who arrives with their own model plugs it in as a first-class citizen. Both paths are equal. Both paths are honored by the same three-component pricing model. The customer's invoice equals value-share on measured savings, plus cost-plus on Stewardly's app and operational expenses, plus provider direct cost when applicable. The whole invoice is capped at the cost-plus-equivalent customer-protection ceiling. The provider direct cost component is path-dependent. On the internal-bundled path Stewardly contracts with the provider on the customer's behalf, so provider direct cost is non-zero and Stewardly applies the cost-plus markup to it. On the bring-your-own path the customer pays the provider directly through their own contract, so the provider direct cost component is zero. Both paths preserve Stewardly's profitability floor through the always-present cost-plus on app and operational expenses component. The platform never runs unprofitably regardless of which path a customer chooses or which stage of the measurement methodology is operating.

The platform is organized around five engines. The five engines are Formational, Relational, Missional, Contextual, and Continuous-Improvement. Formational covers how people learn and grow. Relational covers how people connect to and serve each other. Missional covers how people apply their work to a purpose. Within Missional there are mission specializations. Wealth is the lead specialization today; coaching, pastoral, teaching, healthcare, and community are also enumerated. Contextual covers how the right thing surfaces at the right moment. Continuous-Improvement covers how the whole system gets better over time.

These five engines are ontologically complete. That means there is no user case that falls outside them. Today's implementation is operationally imperfect, and the Continuous-Improvement engine itself drives the closing of those operational gaps over time.

A second completeness commitment is technological. The architecture commits to filling the industry gap of dynamically supporting the universe of current and future artificial-intelligence capabilities and surfaces. The universe of surfaces explicitly includes chat, agents, tools, browsers, deep research, document studios, music studios, video generators, web-app builders, mobile project builders, voice, vision, embedding pipelines, retrieval pipelines, fine-tuning, evaluations, and observability. Each surface is reachable through the same plug-and-play affordances. Adding a new surface or a new provider does not require architectural surgery. This is the technology-side parallel to the engines' ontological completeness. The completeness claim is asymptotic. The platform commits to being able to support whatever the universe contains tomorrow, not to having shipped every surface today.

Provider neutrality is part of this completeness commitment. Bring-your-own is dynamic across the universe of current and future providers. Manus is treated as one provider among many on the bring-your-own side. The bundled out-of-box experience is internal Stewardly orchestration. It is not Manus-as-default. Adding Manus, Anthropic, Open-A-I, Google, an enterprise's self-hosted model, or a future provider follows the same path. The substrate's classifier-driven cost minimization considers all available providers symmetrically. The cost-plus markup on provider direct cost, where applicable, is identical across providers. Provider documentation, onboarding, and create-update-personalize user experiences treat all providers as peers. Audit-log entries identify the provider used for any given intent without privileging any.

Every user at every layer of the administrative hierarchy can create, read, update, delete, and personalize engine applet instances within the controls and constraints set by the layer above them. The hierarchy is five structural layers, plus an orthogonal delegate-scope dimension. The five structural layers are: Tenant, Org or Firm, Team or Supervisor, Professional, and Client. Each layer is anchored to a structural unit. None is defined relative to its relationship to another role. The orthogonal delegate-scope dimension is a property any principal at any layer can use to grant a strictly-subset permission envelope to a delegate. A Professional can delegate to an Assistant. A

Tenant administrator can delegate to a deputy. A Firm administrator can delegate to a deputy. A Team supervisor can delegate to an acting supervisor. A Client can delegate to an authorized representative. The assistant pattern is a delegate-scope of Professional, not a sixth structural layer. The audit log identifies actions taken by a delegate as having been performed under the principal's authority. The create-update-personalize affordance applies at each of the five layers, AND through delegate-scopes. A delegate inherits a strict subset of the principal's affordances and cannot escalate beyond what the principal itself can do. The principal can revoke the delegation at any time. A Tenant administrator constrains what a Firm administrator can configure. A Firm administrator constrains a Team supervisor. A Team supervisor constrains a Professional. A Professional configures their own personalizations and constrains the delegated subset they grant their Assistant. A Client receives what their Professional exposes to them, and may delegate to an authorized representative within the scope the Professional permits. This create-update-personalize affordance at every layer, AND through delegate-scopes, is what makes the infrastructure usable as infrastructure.

The infrastructure framing resolves into three orthogonal axes of dynamism that the architecture supports concurrently. The first axis is horizontal. The platform accommodates the entire universe of bring-your-own artificial-intelligence providers, agents, and related tooling. New providers can be added without architectural surgery. User accessibility and ease of use are first-order requirements of this axis. The second axis is vertical. The three-component pricing model and the two-stage measurement methodology apply across both the bring-your-own path and the bundled-internal path simultaneously. Both use the same measurement methodology. Both preserve Stewardly's profitability floor through the always-present cost-plus on app and operational expenses component. Both honor the same customer protection ceiling. The third axis is extensibility. The platform extensibly accounts for persistent surfaces, persistent artifacts, persistent files, and persistent tools. The Hub is positioned as the persistent-surface front door. The Home, also called Task-Chat, is positioned as the conversational front door. Both front doors are first-class peers anchored to the same persistent substrate.

A crucial strategic posture follows from these axes. The platform deliberately does not position to become its own foundation-model provider. It does not aim to displace Anthropic, Open-A-I, or any equivalent. The non-competition posture has two anchors.

The first anchor is economic. The three-component pricing model is structurally incompatible with foundation-model economics. Foundation-model economics require capturing margin that scales with token volume on differentiated proprietary models. Stewardly's three components scale very differently. Cost-plus on app and operational expenses scales with infrastructure cost, which is sub-linear in token volume because it is amortized across customers. Value-share scales with measured outcome, which is decoupled from token volume by construction. Provider direct cost on the bring-your-own path is zero. None of the three components delivers the token-volume scaling margin that would justify foundation-model competitive investment. On the internal-bundled path the cost-plus markup applied to provider direct cost actually rewards Stewardly for minimizing provider direct cost rather than maximizing it. On the bring-your-own path value-share-only pricing aligns Stewardly's economic interest with the user's chosen provider's success, because more capable providers drive higher measured savings, which drives higher value-share. Either path makes becoming an Anthropic or Open-A-I economically self-defeating. The customer protection ceiling forecloses any degenerate accounting that would attempt it.

The second anchor is compliance. The artificial-intelligence non-compete covenant commits Stewardly contractually not to use user-brought interactions, prompts, outputs, or workflow patterns to train competing models, derive competitive insights, or otherwise compete with the foundation-model providers whose capabilities the platform hosts.

Both anchors are needed. The provider direct cost zero-markup default is in code today. The rest of the economic anchor, which is the cost-plus on app and operational expenses component and the value-share component and the customer protection ceiling, is unbuilt. The compliance anchor is policy intent only and needs an executable invariant plus a counsel-reviewed contractual artifact.

What I found

I audited five repositories on GitHub and three live deployment surfaces. I extracted the deployed bundle from `stewardly.manus.space` and inventoried one hundred and seventy-four lazy-loaded page chunks plus two hundred and twenty-four single-page-application routes. I cloned and inspected the source repositories. I mapped each capability back to its source.

The deployed v3 build at `stewardly.manus.space` is far more complete than the GitHub snapshot in the `mwpenn94/stewardly` repository. The deployed surface ships the full eighteen-tool wealth specialization, fifteen administrative pages, the Manus-

Next foundation for chat and agents and document studio and music studio and video generator and web-app builder and mobile projects, and per-role settings pages for organizations and professionals and managers. Three of the five engine applets are visible in the deployed bundle: Missional, Relational, and Continuous-Improvement. Two are not: Formational and Contextual.

The legacy repository `mwpenn94/stewardly-ai` is the source of most of the deployed v3's depth. The four wealth sub-engines, called UWE, BIE, HE, and SCUI, were ported from `stewardly-ai` into v3's Missional-Wealth path. The relational and customer-relationship-management integrations with GoHighLevel, Dripify, LinkedIn, SMS-iT, and Workable also originated in `stewardly-ai`. Most of these integrations are not yet evidenced in the v3 deployed bundle, which means they are re-incorporation candidates.

The `mwpenn94/emba_modules` repository contains thirty-one production-quality formational tools. These are the AI Quiz, Formula Lab, Exam Simulator, Case Simulator, Hands-Free Study, Study Groups, Peer Groups, Study Buddy, Playlists, Achievements, FS Toolkit, Connection Map, Discipline Deep Dive, Leaderboard, and Progress Export. The deployed v3 bundle does not ship these. They are the natural content for the missing Formational applet.

The `mwpenn94/stewardly-command-center` repository is a property-management product. It does not share a domain with v3. Two of its frontend patterns are domain-portable. One is the kanban contact pipeline with drag and drop between stages. The other is the visual email and outreach workflow builder.

The Manus-Next baseline at `manusnext-mlromfub.manus.space` shows twenty-two routes. All twenty-two are preserved in the deployed v3. There is no app-development capability regression at the route level. Functional verification is still needed.

What is missing

The audit produced a commitment matrix with twenty-seven architectural commitments. The findings cluster into four buckets.

The first bucket is foundational gaps. The three-component pricing formula, the two-stage cost-and-savings measurement methodology, the customer protection ceiling, the bring-your-own first-class path, the dual-anchor non-competition enforcement, the plug-and-play affordance layer for the universe of surfaces and providers, and the provider-neutral abstraction layer that would make Manus one provider among many rather than the de facto default by inheritance from the foundation — none of these yet exist in code. Without them, the infrastructure-positioning claim cannot stand, the strategic non-competition posture is unenforceable, and the technological-completeness and provider-neutrality claims are aspirational rather than architectural. These are foundational, not optional, and they have been promoted to the first wave of Phase Two.

The second bucket is architectural integrity gaps. There is no lint rule preventing engines from importing other engines. There is no executable test for the zero-percent-markup invariant. There is no runtime block-list for the Quantic terms-of-service boundary. The Formational and Contextual engine applet user interfaces are missing. The Hub and Task-Chat shared-substrate parity is not wired and the conversational-to-persistent round-trip is not measured. These cluster in the second wave.

The third bucket is governance gaps. The five-structural-layer permission hierarchy plus the orthogonal delegate-scope primitive, the administrative spectrum runtime gating, the seven counsel-review flag systems, and the wiring of create-update-personalize affordances at every layer and through delegate-scopes are partial or absent. These cluster in waves three and four.

The fourth bucket is re-incorporation opportunities. Formational depth from `emba_modules`, relational depth and integrations from `stewardly-ai`, and portable patterns from `stewardly-command-center` should be brought into v3's architecture in wave seven.

There is also one housekeeping item. The repository root currently contains roughly eighty analysis markdown files. The mandate explicitly limits this to one. That cleanup happens before any other work.

What I propose to do

The work is organized into eight waves. Wave A is foundation. Wave B is architectural integrity. Wave C is identity and governance. Wave D is memory and portability. Wave E is customer-facing language. Wave F is functional verification. Wave G is re-incorporation from legacy sources. Wave H is final convergence.

The total wall-time at agent-tempo is several days of continuous work. This fits within the away-from-keyboard time-box if I am not interrupted mid-wave.

Thirteen judgment calls are documented in the full brief. Each names the chosen approach, the reason, and the trade-off that would cause me to surface it back to you for revision. The most consequential of these is the three-component pricing reconciliation. The customer's invoice equals value-share on measured savings, plus cost-plus on Stewardly's app and operational expenses, plus provider direct cost when applicable on the internal-bundled path. The whole invoice is capped at the cost-plus-equivalent customer-protection ceiling. This formulation reconciles the away-from-keyboard mandate's unified pricing language with your bifurcation message clarifying that internal-app and bring-your-own paths price differently in practice, with your profitability-floor message clarifying that cost-plus must always cover Stewardly's app and operational expenses so the platform never runs unprofitably. The reconciliation is recorded as judgment call D-M in the full brief and is built into Wave A as the unified pricing module before any other pricing logic.

What I need from you

I need confirmation before Phase Two begins. The full brief originally contained seven open questions. One of those questions, the GitHub-versus-deployed-bundle synchronization question, has been permanently closed by user directive. The deployed working tree is to be ported into the `mwpenn94/stewardly` GitHub repository, GitHub becomes the single canonical source of truth from this point forward, all progress is continuously merged into the repo as it happens, and the live web-development project deploys *from* the repo. This is captured as Commitment Twenty-Eight: GitHub as Permanent Canonical Source of Truth. Wave Zero's first task is to port the deployed working tree into the repo as the new authoritative baseline. After the port, every subsequent edit lands in the repo first, then deploys. Six open questions remain in the full brief, none of which block Wave Zero from beginning. The hierarchy structure question, which was raised during this conversation, has also been resolved: the locked answer is Reconciliation C, the five structural layers plus the orthogonal delegate-scope dimension. This is documented as Judgment Call D-N in the full brief.

If you do not answer the six remaining open questions, I will proceed with the documented defaults in the full brief.

To confirm, you can reply with one of three responses. Confirm and proceed. Confirm with changes and list the changes. Or pause and hold.

That is the executive summary. The full brief is available as `PHASE1_BRIEF.md` in the repository.